

Lab3

Dynamic Routing and Network Address Translation

Date : 2021/03/16

Deadline : 2021/03/29 23:59



Outline

- Objective
- Quagga
- Dynamic Routing
- iptables overview
- NAT scenarios
- Lab requirement
- Appendix



Objective

- Dynamic routing configuration
- To learn how Linux kernel handles received packets
- Configure NAT rules on routers with iptables
- Observe packets before/after NAT




Outline

- Objective
- Quagga
 - Introduction
 - Installation
- Dynamic Routing
- iptables overview
- NAT scenarios setup
- Lab requirement
- Appendix



Introduction of Quagga

- Quagga is an open source software that provides routing services
 - Supports common routing protocols: BGP, OSPF, RIP, and IS-IS
 - Consists of a **core daemon Zebra** and separate **routing protocol** daemons
- Routing Protocols (daemons) communicate their best routes to Zebra
- Zebra computes best routes and modifies **kernel routing table** through netlink





Outline

- Objective
- Quagga
 - Introduction
 - Installation
- Dynamic Routing
- iptables overview
- NAT scenarios setup
- Lab requirement
- Appendix



Makefile Dependency Installation

- gawk

- Patten scanning and processing language

```
bash$ sudo apt install gawk -y
```

- libreadline

- Readline and history libraries

```
bash$ sudo apt install libreadline7 libreadline-dev -y
```

- pkg-config

- program used to retrieve information about installed libraries

```
bash$ sudo apt install pkg-config -y
```



Quagga Dependency Installation

- c-ares:
 - Library for asynchronous DNS request
- Download c-ares-1.17.1.tar.gz from e3
- Install c-ares

```
~/Downloads$ tar -xzf c-ares-1.17.1.tar.gz           #unzip this package
~/Downloads$ cd c-ares-1.17.1
~/Downloads/c-ares-1.17.1$ sudo ./configure
#check dependency and generate makefile
~/Downloads/c-ares-1.17.1 $ sudo make                #compile c-ares source code
~/Downloads/c-ares-1.17.1 $ sudo make install        #install c-ares
```




Quagga Installation

- Download quagga-1.2.4.tar.gz from e3
- Install Quagga

```
~/Downloads$ tar -xzf quagga-1.2.4.tar.gz #unzip this package
```

```
~/Downloads$ cd quagga-1.2.4
```

```
~/Downloads/quagga-1.2.4$ sudo ./configure --enable-vtysh --enable-  
user=root --enable-group=root --enable-vty-group=root  
#check dependency and generate makefile with options
```

```
~/Downloads/quagga-1.2.4$ sudo make #compile source code
```

```
~/Downloads/quagga-1.2.4$ sudo make install #install quagga
```

- Copy libzebra to /lib

```
bash $ sudo cp /usr/local/lib/libzebra.so.1 /lib #copy libzebra to /lib
```



Check Quagga Daemons Version

■ Check bgpd version

```
bash$ sudo bgpd -v
```

```
jln@ubuntu:~/Desktop/BGP$ sudo bgpd -v
bgpd version 1.2.4
Copyright 1996-2005 Kunihiro Ishiguro, et al.
configured with:
    --enable-vtysh --enable-user=root --enable-group=root --enable-vty-group
=root
```

■ Check zebra version

```
bash$ sudo zebra -v
```

```
jln@ubuntu:~/Desktop/BGP$ sudo zebra -v
zebra version 1.2.4
Copyright 1996-2005 Kunihiro Ishiguro, et al.
configured with:
    --enable-vtysh --enable-user=root --enable-group=root --enable-vty-group
=root
```




Outline

- Objective
- Quagga
- **Dynamic Routing**
- iptables overview
- NAT scenarios setup
- Lab requirement
- Appendix



Example

- Create two ASs in mininet
 - AS: Autonomous System




- Routers use BGP to exchange routes
 - BGP: Border Gateway Protocol



Network Configuration

- Both r1 and r2 run BGP and Zebra daemons
 - Each router needs a configuration file for each daemon
- Create configuration files for daemons





Python script for example scenario

- Download example.py from e3
- Create and edit configuration files for daemons on r1 and r2
- Put configuration files in the directory **specified by example.py**
 - E.g., configuration files and example.py in the same directory

```
jln@ubuntu:~/Desktop/exampleScen$ tree
.
├── bgp_r1.conf
├── bgp_r2.conf
├── example.py
└── zebra.conf
```



Run BGP and Zebra daemons on Routers

- Create a directory for pid-files of daemons

```
bash$ sudo mkdir /var/run/quagga/
```

- Python script that runs Zebra and bgpd daemons on r1 and r2

```
58 r1.cmd('zebra -f ./zebra.conf -d -i /var/run/quagga/zebraR1.pid')
59 r1.cmd('bgpd -f ./bgp_r1.conf -d -i /var/run/quagga/bgpdR1.pid')
60 r2.cmd('zebra -f ./zebra.conf -d -i /var/run/quagga/zebraR2.pid')
61 r2.cmd('bgpd -f ./bgp_r2.conf -d -i /var/run/quagga/bgpdR2.pid')
```

- -f: specify a config file
- -d: runs in daemon mode
- -i: create a pid-file for this daemon



Configuration for Zebra Daemons

- Define hostname and password for Zebra daemon
 - For telnet to Zebra daemon

```
! Configuration for zebra (Note: it is the same for all routers)
!  
hostname zebra  
password nscap  
log stdout  
!
```




Configuration for BGP Daemon on r1

! BGP configuration for r1

```
!
hostname r1
password nscap
!
router bgp 65000
  bgp router-id 10.0.0.3
  timers bgp 3 9
  neighbor 10.0.0.4 remote-as 65001
  neighbor 10.0.0.4 ebgp-multihop
  neighbor 10.0.0.4 timers connect 5
  neighbor 10.0.0.4 advertisement-interval 5
  network 192.168.1.0/24
!
log stdout
```

ASN 65000

ASN 65001

192.168.1.0 /24



h1

10.0.0.3/8



r1

10.0.0.4/8



r2

192.168.2.0/24



h2

ASN 65000 —

ASN 65001 —



Configuration for BGP Daemon on r2

```
! BGP configuration for r2
```

```
!
```

```
hostname r2
```

```
password nscap
```

```
!
```

```
router bgp 65001
```

```
bgp router-id 10.0.0.4
```

```
timers bgp 3 9
```

```
neighbor 10.0.0.3 remote-as 65000
```

```
neighbor 10.0.0.3 ebgp-multihop
```

```
neighbor 10.0.0.3 timers connect 5
```

```
neighbor 10.0.0.3 advertisement-interval 5
```

```
network 192.168.2.0/24
```

```
!
```

```
log stdout
```

ASN 65000

192.168.1.0 /24



h1



10.0.0.3/8

r1

10.0.0.4/8



r2



h2

ASN 65001

192.168.2.0/24

ASN 65000 —

ASN 65001 —



Check Route in routing table

■ Check Route

```
mininet> r1 route
```

```
mininet> r1 route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0          255.0.0.0        U        0      0        0 r1-eth0
192.168.1.0        0.0.0.0          255.255.255.0    U        0      0        0 r1-eth1
192.168.2.0        10.0.0.4          255.255.255.0    UG       20     0        0 r1-eth0
```

```
mininet> r2 route
```

```
mininet> r2 route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0          255.0.0.0        U        0      0        0 r2-eth0
192.168.1.0        10.0.0.3          255.255.255.0    UG       20     0        0 r2-eth0
192.168.2.0        0.0.0.0          255.255.255.0    U        0      0        0 r2-eth1
```



Check Route in Zebra

- Telnet r1 zebra daemons (on port 2601)

```
mininet> r1 xterm & #invoke a terminal for r1  
r1> telnet 127.0.0.1 2601
```

```
User Access Verification
```

```
Password:  
zebra> █
```

Type in password defined in zebra.conf

- Show bgp route in r1

```
zebra> show ip route bgp
```

```
zebra> show ip route bgp  
Codes: K - kernel route, C - connected, S - static, R - RIP,  
        O - OSPF, I - IS-IS, B - BGP, P - PIM, A - Babel, N - NHRP,  
        > - selected route, * - FIB route
```

```
B>* 192.168.2.0/24 [20/0] via 10.0.0.4, r1-eth0, 00:18:25  
zebra> █
```



Check route in bgpd

- Telnet r1 bgpd daemons (om port 2605)

```
r1> telnet 127.0.0.1 2605
```

```
User Access Verification
```

```
Password:
```

```
r1> █
```

Type in Password defined in bgp_r1.conf

- Show r1 bgp neighbor summary

```
zebra> show ip bgp summary
```

```
Total number of routes received 1  
r1> show ip bgp summary  
BGP router identifier 10.0.0.3, local AS number 65000  
RIB entries 3, using 336 bytes of memory  
Peers 1, using 9088 bytes of memory
```

| Neighbor | V | AS | MsgRcvd | MsgSent | TblVer | InQ | OutQ | Up/Down | State/PfxRcd |
|----------|---|-------|---------|---------|--------|-----|------|----------|--------------|
| 10.0.0.4 | 4 | 65001 | 426 | 429 | 0 | 0 | 0 | 00:21:12 | 1 |

```
Total number of neighbors 1
```

```
Total num. Established sessions 1  
Total num. of routes received 1
```



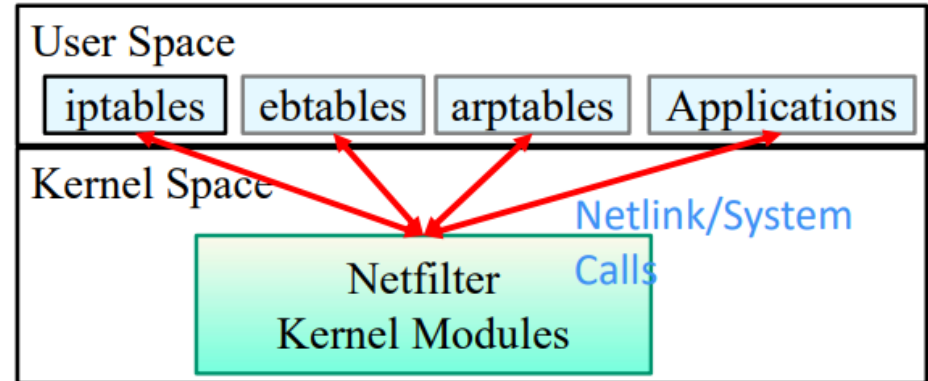
Outline

- Objective
- Quagga
- Dynamic Routing
- **iptables**
 - Overview
 - Basic usage
- NAT scenarios
- Lab requirement
- Appendix



iptables overview

- A user-space utility program for configuring IP packet filter rules of Linux kernel firewall
 - Linux kernel firewall implemented as different **Netfilter modules**.
- Netfilter: a framework inside the Linux kernel that allows kernel modules to register **callback** functions at different locations (hooks) of the Linux network stack.
 - – A **registered callback** function is called back for every packet that traverses the respective hook within the Linux network stack.





Component of iptables

- **Tables:** files that join similar actions.
 - Contains a number of built-in chains or user-defined chains.
- **Chains:** a list of **rules** which can match a set of packets
 - When receives a packet, iptables finds the appropriate table;
 - Then apply the chain of **rules** on the packet until it finds a match.
- **Rules:** specifies what to do with a packet that matches.
 - can block one type of packet, or
 - forward another type of packet.
- **Targets:** a decision of what to do with a packet.
 - Typically, Accept, Drop, or Reject (which sends an error back to the sender)

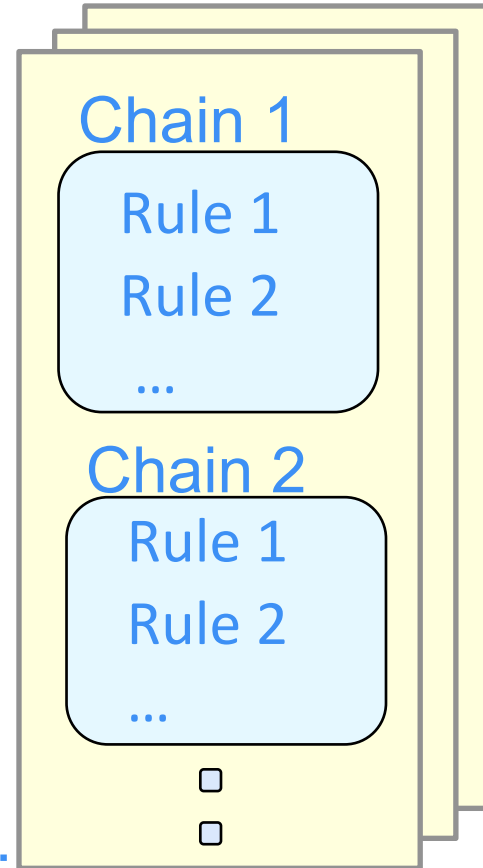


Table filter/...



Tables and Chains

■ Tables

- filter: packet filtering, default table
- nat: NAT operation
- mangle: add tag on packet (for QoS or load distribution)
- raw: mainly for exemptions from connection tracking

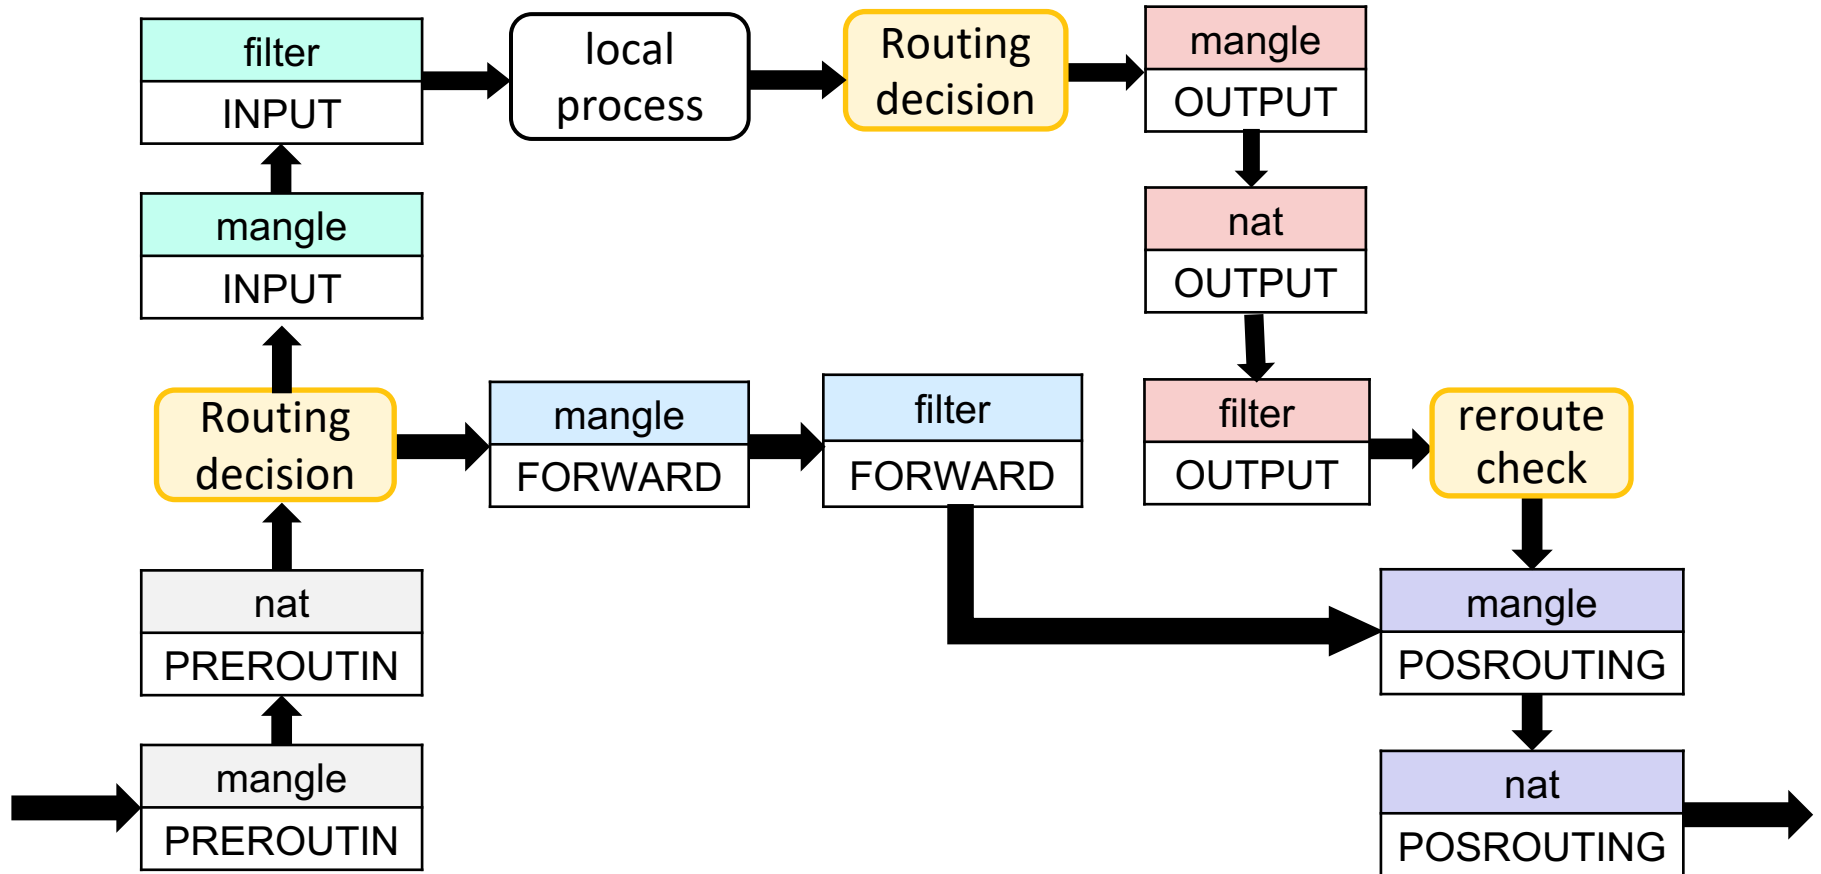
■ Five Predefined Chains (mapping to the five available Netfilter hooks)

- **PREROUTING**: for packets before a routing decision is made.
- **INPUT**: for packets destined to local sockets
- **FORWARD**: for packets being routed through the machine.
- **OUTPUT**: for locally-generated packets.
- **POSTROUTING**: for packets about to go out after Routing decision has been made.





Simplified Netfilter Network Layer Packet Flow





Outline

- Objective
- Quagga
- Dynamic Routing
- **iptables**
 - Overview
 - Basic Operations
- NAT scenarios
- Lab requirement
- Appendix



iptables – Chains and Rules

- Show all chains and rules in a table

```
bash$ sudo iptables -nvL -t [table name]
```

- n: number
- v: detail
- L: List

```
jin@ubuntu:~/Desktop$ sudo iptables -nvL -t nat
Chain PREROUTING (policy ACCEPT 6586 packets, 924K bytes)
pkts bytes target      prot opt in      out     source      destination
  4    204 DOCKER      all  --  *       *       0.0.0.0/0   0.0.0.0/0

Chain INPUT (policy ACCEPT 6552 packets, 921K bytes)
pkts bytes target      prot opt in      out     source      destination

Chain OUTPUT (policy ACCEPT 40343 packets, 3435K bytes)
pkts bytes target      prot opt in      out     source      destination
  0     0 DOCKER      all  --  *       *       0.0.0.0/0   !127.0.0.0/8

Chain POSTROUTING (policy ACCEPT 40343 packets, 3435K bytes)
pkts bytes target      prot opt in      out     source      destination
  0     0 MASQUERADE  all  --  *       !docker0 172.17.0.0/16 0.0.0.0/0

Chain DOCKER (2 references)
pkts bytes target      prot opt in      out     source      destination
  0     0 RETURN     all  --  docker0 *       0.0.0.0/0   0.0.0.0/0
```

Default Chains

Custom Chains

Default Policy



iptables – Matching Fields and Actions

■ Matching fields

```
pkts bytes target prot opt in out source destination
0 0 DOCKER all -- * * 0.0.0.0/0 0.0.0.0/0 ADDRTYPE match dst-type LOCAL
```

■ Actions

■ Target indicate

- Actions of rules like DNAT/ SNAT/ MASQUERADE
- Jump to another chains

```
Chain POSTROUTING (policy ACCEPT 38271 packets, 3243K bytes)
pkts bytes target prot opt in out source destination
0 0 MASQUERADE all -- * !docker0 172.17.0.0/16 0.0.0.0/0
```

```
pkts bytes target prot opt in out source destination
0 0 DOCKER all -- * * 0.0.0.0/0 0.0.0.0/0 ADDRTYPE match dst-type LOCAL
```



Iptables – Adding Rules

- Add rules to a chain of a table

```
bash$ sudo iptables -t [table] -A [chain] [match field] -j [Actions]
```

- -A: append
- -I: insert

- Example:

```
bash$ sudo iptables -t nat -A POSTROUTING -s 172.20.0.0/16 -d 172.87.0.0/16 -o eth2 -j SNAT --to-source 140.113.194.239
```

- Append a SNAT rules chain POSTROUTING of NAT table if packet matched following fields
 - source address is in 172.20.0.0/16
 - destination address is in 172.87.0.0/16
 - the output interface is eth2

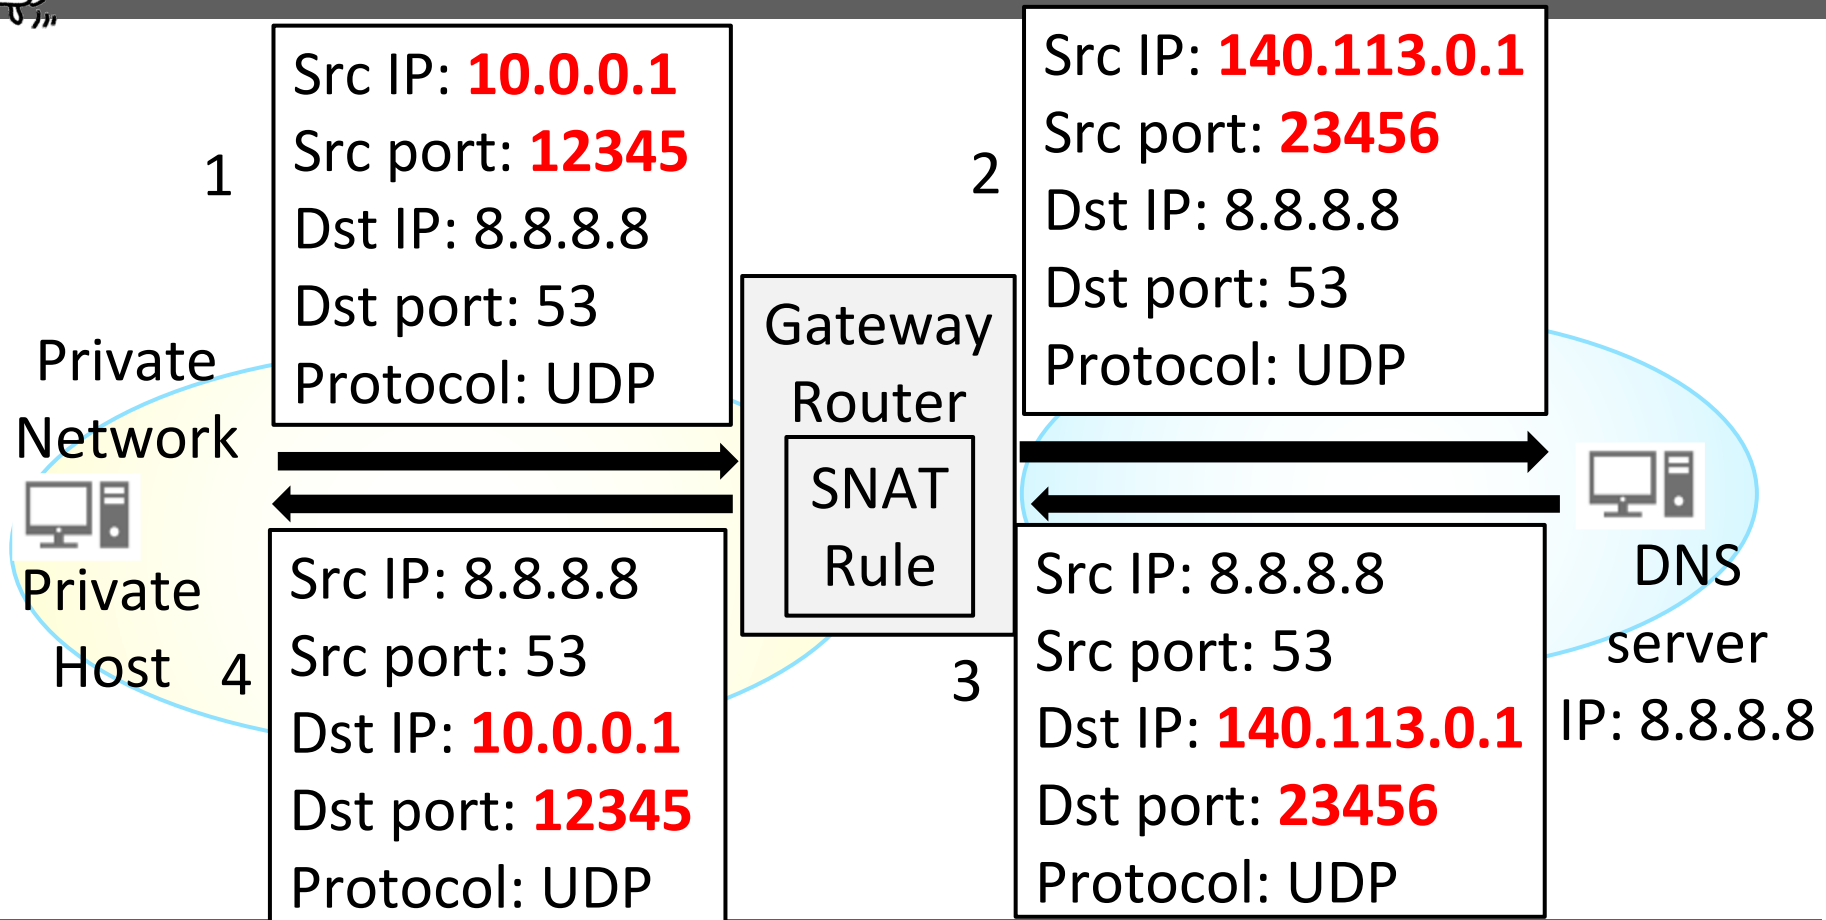


Outline

- Objective
- Quagga
- Dynamic Routing
- iptables overview
- **NAT scenarios**
 - Source NAT
 - Destination NAT
- Lab requirement
- Appendix

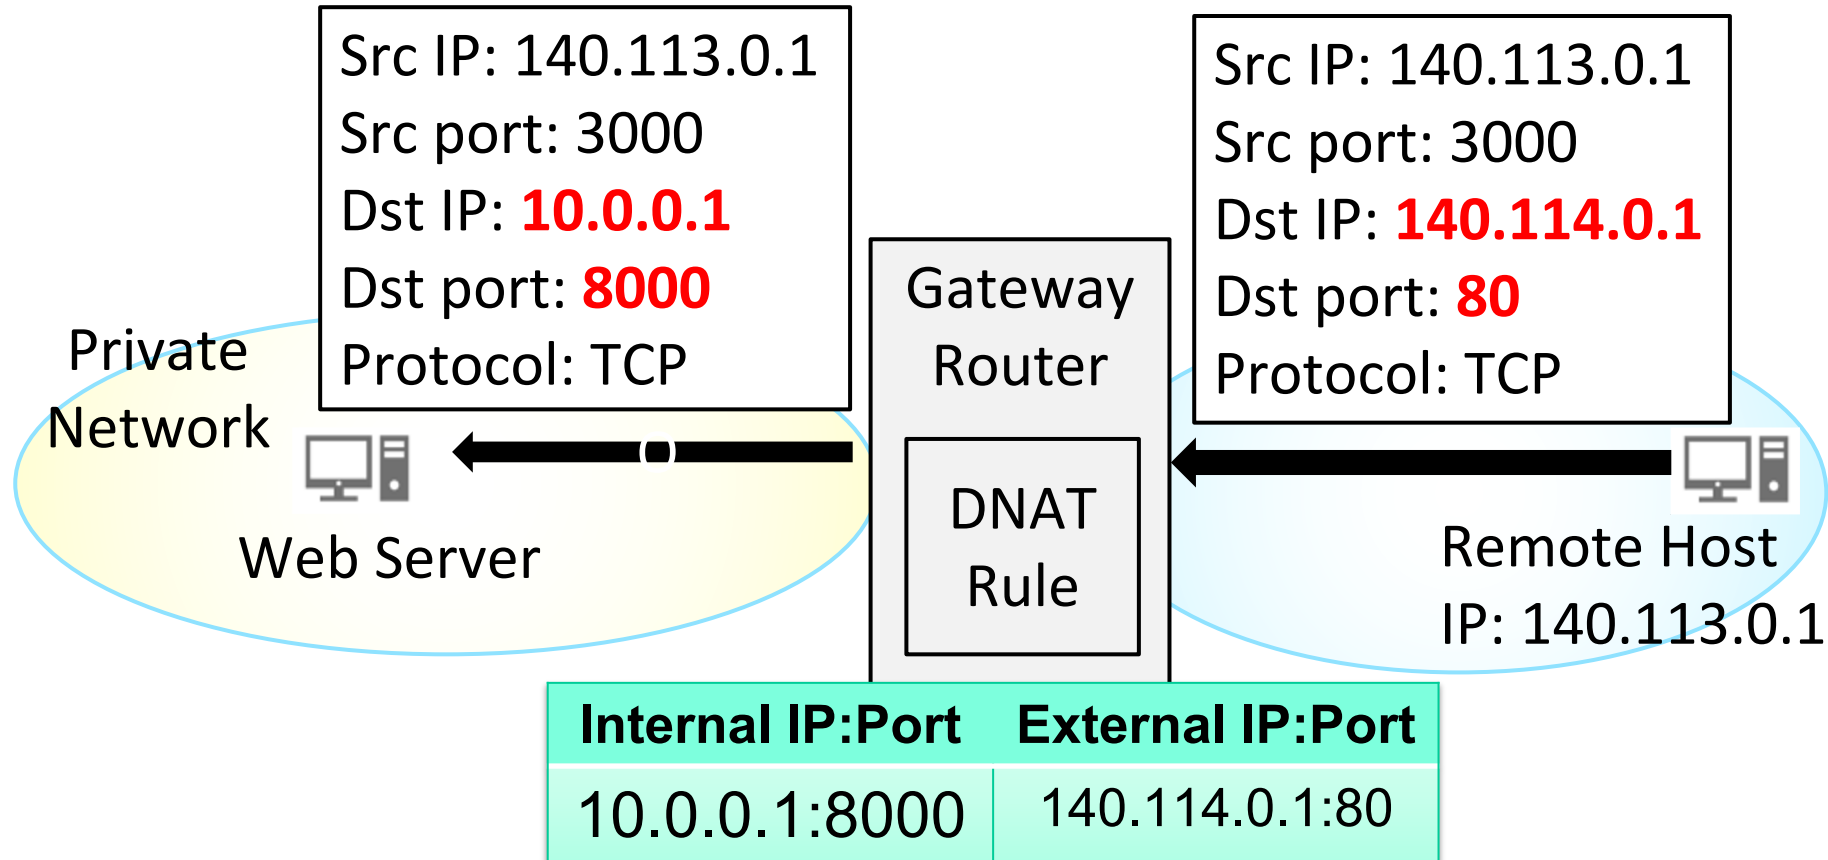


Source NAT





Destination NAT



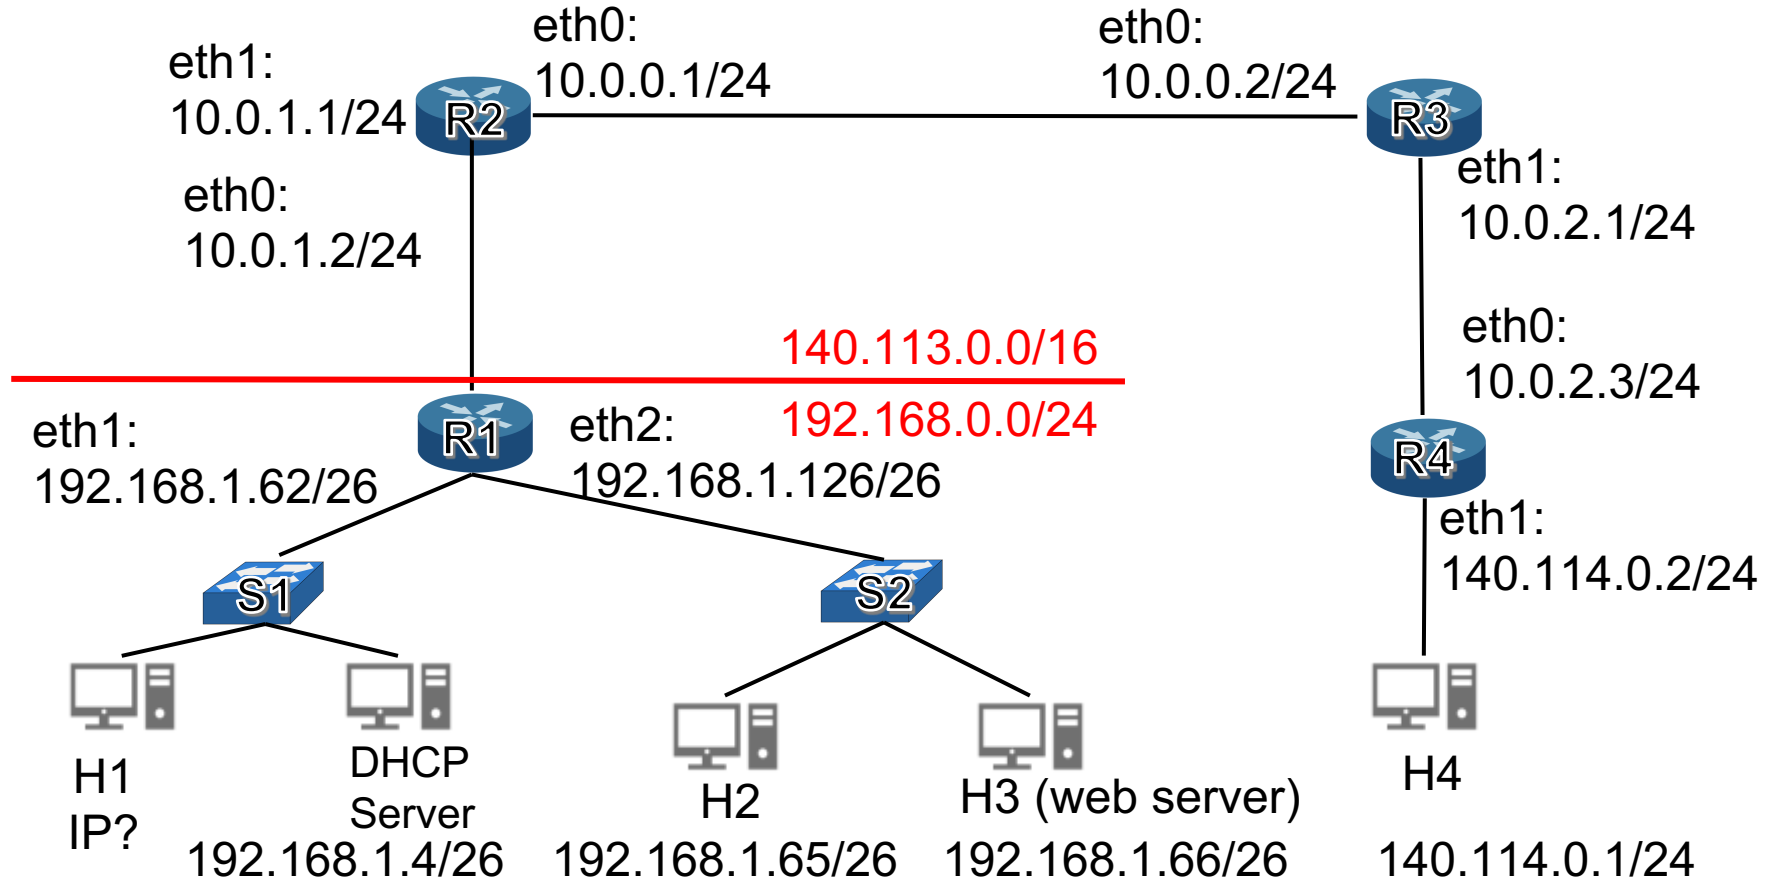


Outline

- Objective
- Quagga
- Dynamic Routing
- iptables overview
- NAT scenarios
- Lab requirement
 - Topology
 - Part1: dynamic routing
 - Part2: configure NAT rules
 - About submit



Lab Topology





Part1: Topology Creation and Configuration

- Download topology.py from e3
- Edit bgp and zebra configuration files for routers
 - bgp_<router>.conf
 - zebra.conf
- Execute topology.py

```
bash$ sudo python topology.py
```

- Routing tables before enabling BGP



```
mininet> r1 route
Kernel IP routing table
Destination    Gateway      Genmask      Flags Metric Ref    Use Iface
10.0.1.0       0.0.0.0      255.255.255.0 U        0      0      0 r1-eth0
192.168.1.0    0.0.0.0      255.255.255.192 U        0      0      0 r1-eth1
192.168.1.64   0.0.0.0      255.255.255.192 U        0      0      0 r1-eth2
mininet> r2 route
Kernel IP routing table
Destination    Gateway      Genmask      Flags Metric Ref    Use Iface
10.0.0.0       0.0.0.0      255.255.255.0 U        0      0      0 r2-eth0
10.0.1.0       0.0.0.0      255.255.255.0 U        0      0      0 r2-eth1
mininet> r3 route
Kernel IP routing table
Destination    Gateway      Genmask      Flags Metric Ref    Use Iface
10.0.0.0       0.0.0.0      255.255.255.0 U        0      0      0 r3-eth0
10.0.2.0       0.0.0.0      255.255.255.0 U        0      0      0 r3-eth1
mininet> r4 route
Kernel IP routing table
Destination    Gateway      Genmask      Flags Metric Ref    Use Iface
10.0.2.0       0.0.0.0      255.255.255.0 U        0      0      0 r4-eth0
140.114.0.0    0.0.0.0      255.255.255.0 U        0      0      0 r4-eth1
mininet>
```



Part1: BGP packet capturing

- Run wireshark on node r2 and r3 to capture BGP packets

```
mininet> r2 wireshark & #listen at r2-eth0
```

```
mininet> r3 wireshark & #listen at r3-eth0
```

```
mininet> r3 wireshark & #listen at r3-eth1
```

- Run BGP and Zebra daemons on every router nodes ([r1-r4])

```
mininet> r1 zebra -f ./configs/zebra.conf -d -i  
/var/run/quagga/zebra1.pid
```

```
mininet> r1 bgpd -f ./configs/bgp_r1.conf -d -i  
/var/run/quagga/bgpd1.pid
```



Part1: Check routing tables after enabling BGP

- Check routing tables again

```
mininet> r1 route
Kernel IP routing table
Destination        Gateway           Genmask          Flags Metric Ref    Use Iface
10.0.1.0           0.0.0.0          255.255.255.0    U        0      0      0 r1-eth0
140.114.0.0        10.0.1.1         255.255.0.0      UG       20     0      0 r1-eth0
192.168.1.0        0.0.0.0          255.255.255.192  U        0      0      0 r1-eth1
192.168.1.64       0.0.0.0          255.255.255.192  U        0      0      0 r1-eth2
mininet> r2 route
Kernel IP routing table
Destination        Gateway           Genmask          Flags Metric Ref    Use Iface
10.0.0.0           0.0.0.0          255.255.255.0    U        0      0      0 r2-eth0
10.0.1.0           0.0.0.0          255.255.255.0    U        0      0      0 r2-eth1
140.113.0.0        10.0.1.2         255.255.0.0      UG       20     0      0 r2-eth1
140.114.0.0        10.0.0.2         255.255.0.0      UG       20     0      0 r2-eth0
mininet> r3 route
Kernel IP routing table
Destination        Gateway           Genmask          Flags Metric Ref    Use Iface
10.0.0.0           0.0.0.0          255.255.255.0    U        0      0      0 r3-eth0
10.0.2.0           0.0.0.0          255.255.255.0    U        0      0      0 r3-eth1
140.113.0.0        10.0.0.1         255.255.0.0      UG       20     0      0 r3-eth0
140.114.0.0        10.0.2.3         255.255.0.0      UG       20     0      0 r3-eth1
mininet> r4 route
Kernel IP routing table
Destination        Gateway           Genmask          Flags Metric Ref    Use Iface
10.0.2.0           0.0.0.0          255.255.255.0    U        0      0      0 r4-eth0
140.113.0.0        10.0.2.1         255.255.0.0      UG       20     0      0 r4-eth0
140.114.0.0        0.0.0.0          255.255.255.0    U        0      0      0 r4-eth1
```



Part1: Question

1. Take routing tables screenshot before/after on [r1-r4] (10%)
2. Telnet zebra and bgpd daemons of [r1-r4] and take screenshots of routes in zebra and bgpd daemons. (10%)
3. Capture BGP packets from wireshark and take screenshot to verify your answer for the following questions (20%)

3-1. Show BGP packets (OPEN, UPDATE, KEEP ALIVE) exchanged by r2 and r3

3-2. What will happen to the routing table if you set r4-eth0 down?

```
mininet> r4 ip link set r4-eth0 down
```

```
mininet> [r1-r4] route #check routing tables
```

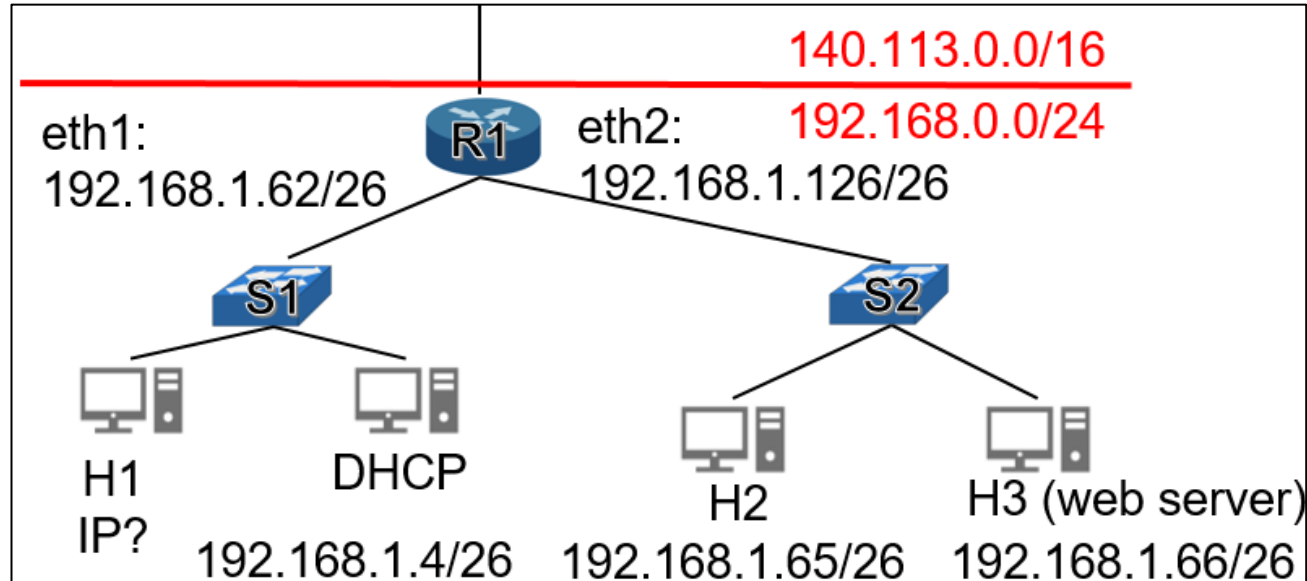
3-3. How does r3 know r4 is unreachable? Explain how

3-4. How does r2 know r4 is unreachable? Explain how



Part2: Source NAT

- Configure r1 to perform Source NAT with iptables rules
 - Use 140.113.0.30 for network 192.168.1.0/26
 - Use 140.113.0.40 for network 192.168.1.64/26





Part2: Destination NAT

- Run Http server at node h3

```
mininet> h3 python -m SimpleHTTPServer &
```

- SimpleHTTPServer

Python build-in script, listen at port 8000 by default


- Configure r1 to perform Destination NAT with iptables rules

- Mapping 140.113.0.40:80 to 192.168.1.66:8000

- h4 sends Http GET to h3

```
mininet> h4 curl 140.113.0.40:80
```

- 1. Take screenshot of curl result (10%)





Part2: Question

2. Check reachability and take screenshot (10%)

```
mininet> h1 ping h4 -c 1  
mininet> h2 ping h4 -c 1  
mininet> h3 ping h4 -c 1
```

3. Run wireshark on r1 to take screenshot of input/output packet (10%)

- Explain the difference of packet headers

```
mininet> r1 wireshark & #listen at r1-eth0  
mininet> r1 wireshark & #listen at r1-eth1  
mininet> r1 wireshark & #listen at r1-eth2  
mininet> h1 ping h4 -c 1  
mininet> h2 ping h4 -c 1
```



Report Submission

■ Files

- <StudentID>_topo.py (10%)
 - With NAT configuration
- bgp_<router>.conf and zebra.conf (20%)
- A Report: lab3_<studentID>.pdf (70%)
 - Screenshot and answers

■ Submission

- Zip all files into a zip file
 - Name: lab3_<studentID>.zip
 - Wrong filename or format will deduct scores (-5%)



THANK
YOU



Appendix

- iptables man page
 - <https://linux.die.net/man/8/iptables>
- Quagga
 - <https://www.quagga.net/docs/quagga.pdf>