

0716306 project1 report

Question 1 : Configuration command

以下的 command 都是在 VM 的 terminal 下達

a. BRG1

```
# start docker BRG1
· docker start BRG1

# create a veth pair between interface BRG1-eth1 and br0-BRG1(BRG1-eth1 is
the interface to br0, br0-BRG1 is the interface to br0)
· sudo ip link add BRG1-eth1 type veth peer name br0-BRG1

# put the interface BRG1-eth1 into BRG1 container(br0 is created in VM)
· sudo ip link set BRG1-eth1 netns $(sudo dockeer inspect -f '{{.State.Pid}}')

BRG1

# bridge the interface br0 with br0-BRG1
· sudo brctl addif br0 br0-BRG1

# set the interface up
· sudo ip link set br0-BRG1 up

# enable ipv4 forwarding of BRG1
· sudo docker exec -it BRG1 sysctl -w net.ipv4.ip_forward=1

# set up BRG1-eth1
· sudo docker exec -it BRG1 ip link set BRG1-eth1 up
```

b. BRGr

```
# enable ipv4 forwarding
· sudo docker exec -it BRGr sysctl -w net.ipv4.ip_forward=1

# set up interfaces
· sudo docker exec -it BRGr ip link set BRGr-eth0 up
· sudo docker exec -it BRGr ip link set BRGr-eth1 up

# assign static ip to BRGr-eth0
· sudo docker exec -it BRGr ifconfig BRGr-eth0 140.113.0.2

# copy GRE tunnel creation program to BRGr(executable file)
· docker cp gre_creation BRGr:/gre_creation

# execute the GRE tunnel creation file
· ./gre_creation

# add route to ISP
· sudo docker exec -it BRGr route add -net 172.27.0.0/16 gw 140.114.0.1
```

c. Edge router

```
# enable ipv4 forwarding
· sudo docker exec -it edge_router sysctl -w net.ipv4.ip_forward=1

# allocate ip for interface and set up
· sudo docker exec -it edge_router ifconfig edge-eth1 140.114.0.1
· sudo docker exec -it edge_router ifconfig edge-eth0 172.27.0.1
· sudo docker exec -it edge_router ip link set edge-eth0 up
· sudo docker exec -it edge_router ip link set edge-eth1 up

# create a dhcpd.conf named dhcpd_edge.conf
· sudo docker exec -it edge_router echo -e "subnet 172.27.0.0 netmask
255.255.255.0 {\n    range 172.27.0.10 172.27.0.250;\n    option routers
172.27.0.1;\n}" > dhcpd_edge.conf

# run DHCP server and allocate IP to interface of BRG1
· sudo docker exec -it edge_router /usr/sbin/dhcpd 4 -cf ./dhcpd_edge.conf
edge-eth0

# change the iptables FORWARD chain setting to MASQUERADE
· sudo docker exec -it edge_router iptables -t nat -A POSTROUTING -s
172.27.0.0/24 -o edge-eth1 -j MASQUERADE
```

```
root@27f28b593be4:/# iptables -t nat -nvL
Chain PREROUTING (policy ACCEPT 4 packets, 824 bytes)
pkts bytes target      prot opt in      out     source      destination
Chain INPUT (policy ACCEPT 4 packets, 824 bytes)
pkts bytes target      prot opt in      out     source      destination
Chain OUTPUT (policy ACCEPT 4 packets, 245 bytes)
pkts bytes target      prot opt in      out     source      destination
Chain POSTROUTING (policy ACCEPT 4 packets, 245 bytes)
pkts bytes target      prot opt in      out     source      destination
0      0 MASQUERADE  all  --  *      edge-eth1  172.27.0.0/24  0.0.0.0/0
```

```
# add route to ISP
· sudo docker exec -it edge_router route add -net 140.113.0.0/16 gw
140.114.0.1
```

d. GWr

```
# enable ipv4 forwarding
· sudo sysctl -w net.ipv4.ip_forward=1

# create a dhcpd.conf named dhcpd_GWr.conf
· sudo echo -e "subnet 20.0.0.0 netmask 255.0.0.0 {\n    range 20.0.0.10
20.0.0.250;\n    option routers 10.0.2.15;\n    option domain-name-
servers 8.8.8.8;\n}" > dhcpd_GWr.conf

# run DHCP server
· sudo /usr/sbin/dhcpd 4 -cf ./dhcpd_GWr.conf GWr-eth0

# set NAT
· sudo iptables -t nat -A POSTROUTING -s 20.0.0.0/8 -o enp0s3 -j
MASQUERADE
```

Chain PREROUTING (policy ACCEPT 25 packets, 12428 bytes)									
pkts	bytes	target	prot	opt	in	out	source	destination	
22	11688	DOCKER	all	--	*	*	0.0.0.0/0	0.0.0.0/0	ADDRTYPE match dst-type LOCAL
Chain INPUT (policy ACCEPT 20 packets, 10536 bytes)									
pkts	bytes	target	prot	opt	in	out	source	destination	
0	0	DOCKER	all	--	*	*	0.0.0.0/0	!127.0.0.0/8	ADDRTYPE match dst-type LOCAL
Chain OUTPUT (policy ACCEPT 946 packets, 78231 bytes)									
pkts	bytes	target	prot	opt	in	out	source	destination	
0	0	MASQUERADE	all	--	*	!docker0	172.17.0.0/16	0.0.0.0/0	
0	0	MASQUERADE	all	--	*	enp0s3	20.0.0.0/8	0.0.0.0/0	
Chain DOCKER (2 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
0	0	RETURN	all	--	docker0	*	0.0.0.0/0	0.0.0.0/0	

Question 2 : Interface information

a. BRG1

```
BRG1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
ether 8a:51:d9:99:eb:87 txqueuelen 1000 (Ethernet)
RX packets 194 bytes 66348 (66.3 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 0 bytes 0 (0.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

BRG1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 172.27.0.5 netmask 255.255.255.0 broadcast 172.27.0.255
ether 7e:97:d1:95:76:a4 txqueuelen 1000 (Ethernet)
RX packets 64 bytes 6959 (6.9 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 10 bytes 1452 (1.4 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
inet 127.0.0.1 netmask 255.0.0.0
loop txqueuelen 1000 (Local Loopback)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 0 bytes 0 (0.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

b. BRG2

```
BRG2-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
ether 32:af:bb:81:cb:74 txqueuelen 1000 (Ethernet)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 0 bytes 0 (0.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
inet 127.0.0.1 netmask 255.0.0.0
loop txqueuelen 1000 (Local Loopback)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 0 bytes 0 (0.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

c. BRGr

```
BRGr-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 140.113.0.2 netmask 255.255.0.0 broadcast 140.113.255.255
ether 3a:9d:4d:de:dc:08 txqueuelen 1000 (Ethernet)
RX packets 6 bytes 476 (476.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 6 bytes 476 (476.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

BRGr-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 20.0.0.10 netmask 255.0.0.0 broadcast 20.255.255.255
ether ea:fe:40:a5:19:1d txqueuelen 1000 (Ethernet)
RX packets 85 bytes 9574 (9.5 KB)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 22 bytes 2516 (2.5 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
ether aa:05:69:59:9a:2f txqueuelen 1000 (Ethernet)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 2 bytes 108 (108.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
inet 127.0.0.1 netmask 255.0.0.0
loop txqueuelen 1000 (Local Loopback)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 0 bytes 0 (0.0 B)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Question 3

Question 4

在 h1 發送封包後，BRGr 會在 MAC table 中記錄 h1 從哪個 gretap 進入，因此在 GWr 要將封包發送出去時，Linux kernel 可以透過 GRE packet 最外層 (outer ethernet header)得知要轉送到 h1 或 h2 的封包應該由哪個 gretap 出去